

Bayesian SIR Model

Rob Deardon (Calgary) and Caitlin Ward (Minnesota)

Writing the SIR model

```
library(nimble)

## nimble version 0.12.2 is loaded.
## For more information on NIMBLE and a User Manual,
## please visit https://R-nimble.org.
##
## Attaching package: 'nimble'
## The following object is masked from 'package:stats':
##
##   simulate
SIR_code <- nimbleCode({

  S[1] <- N - I0 - R0
  I[1] <- I0
  R[1] <- R0

  probIR <- 1 - exp(-gamma)

  ### loop over time
  for(t in 1:tau) {

    probSI[t] <- 1 - exp(- beta * I[t] / N)

    Istar[t] ~ dbin(probSI[t], S[t])
    Rstar[t] ~ dbin(probIR, I[t])

    # update S, I, R
    S[t + 1] <- S[t] - Istar[t]
    I[t + 1] <- I[t] + Istar[t] - Rstar[t]
    R[t + 1] <- R[t] + Rstar[t]

  }

  # priors
  beta ~ dgamma(1, 1)
  gamma ~ dgamma(aa, bb)

})
```

Simulating epidemics

Here we specify the population size $N = 10,000$, 5 initially infectious individuals, and simulate 100 days of the epidemic.

```
constantsList <- list(N = 10000,
                    I0 = 5,
                    R0 = 0,
                    tau = 100)

sirModel <- nimbleModel(SIR_code,
                      constants = constantsList)

## Defining model
## Building model
## Running calculate on model
## [Note] Any error reports that follow may simply reflect missing values in model variables.
## Checking model sizes and dimensions
## [Note] This model is not fully initialized. This is not an error.
## To see which variables are not initialized, use model$initializeInfo().
## For more information on model initialization, see help(modelInitialization).

# exclude data from parent nodes
dataNodes <- c('Istar', 'Rstar')
dataNodes <- sirModel$expandNodeNames(dataNodes, returnScalarComponents = TRUE)
parentNodes <- sirModel$getParents(dataNodes, stochOnly = TRUE)
parentNodes <- parentNodes[-which(parentNodes %in% dataNodes)]
parentNodes <- sirModel$expandNodeNames(parentNodes, returnScalarComponents = TRUE)
nodesToSim <- sirModel$getDependencies(parentNodes, self = FALSE, downstream = T)
```

We can simulate using various values of β and γ to specify various reproductive numbers.

In all simulations the mean infectious period is 5 days.

```
pal <- c('forestgreen', 'red', 'blue')

par(mfrow = c(2,2))

# simulation 1
initsList <- list(beta = 0.8,
                 gamma = 0.2)
sirModel$setInits(initsList)

set.seed(1)
sirModel$simulate(nodesToSim, includeData = TRUE)

plot(sirModel$S, type = 'l', col = pal[1], ylim = c(0, 1.3e4),
     main = paste0('R0 = ', sirModel$beta / sirModel$gamma), lwd = 2, ylab = "Population Count")
lines(sirModel$I, col = pal[2], lwd = 2)
lines(sirModel$R, col = pal[3], lwd = 2)
legend('topright', c('S', 'I', 'R'), col = pal, lwd = 2,
      bty = 'n', horiz = T)
```

```

plot(sirModel$I, type = 'l', col = pal[1], ylim = c(0, 5000),
     main = paste0('R0 = ', sirModel$beta / sirModel$gamma), lwd = 2, ylab = "Population Count")

# simulation 2
initsList <- list(beta = 0.4,
                  gamma = 0.2)
sirModel$setInits(initsList)

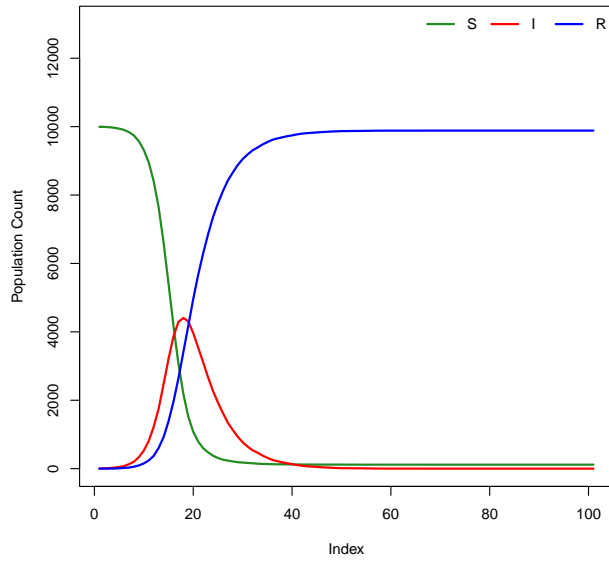
set.seed(1)
sirModel$simulate(nodesToSim, includeData = TRUE)

plot(sirModel$S, type = 'l', col = pal[1], ylim = c(0, 10000),
     main = paste0('R0 = ', sirModel$beta / sirModel$gamma), lwd = 2, ylab = "Population Count")
lines(sirModel$I, col = pal[2], lwd = 2)
lines(sirModel$R, col = pal[3], lwd = 2)
legend('topright', c('S', 'I', 'R'), col = pal, lwd = 2,
      bty = 'n', horiz = T)

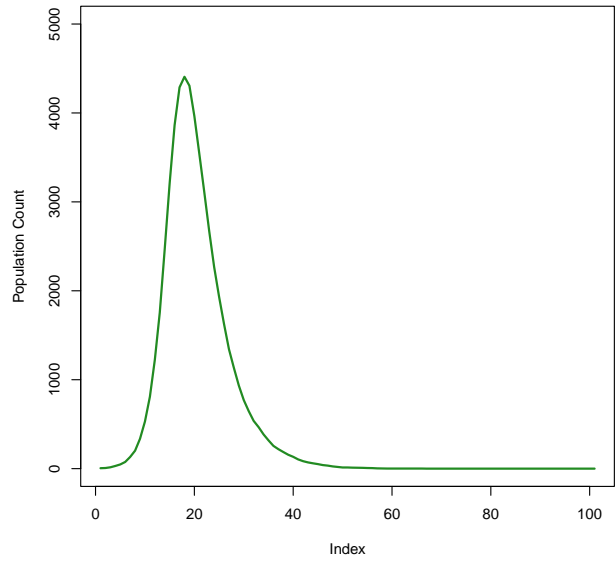
plot(sirModel$I, type = 'l', col = pal[1], ylim = c(0, 2000),
     main = paste0('R0 = ', sirModel$beta / sirModel$gamma), lwd = 2)

```

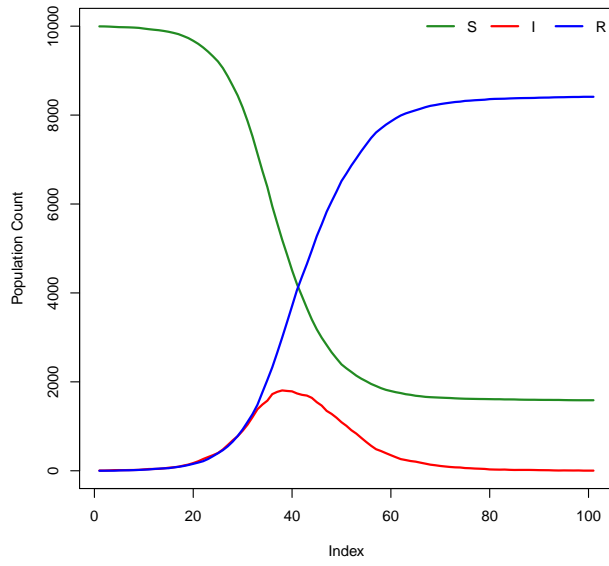
R0 = 4



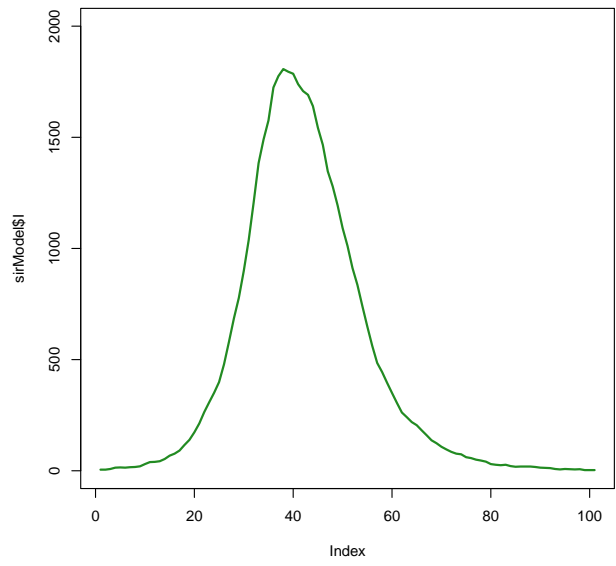
R0 = 4



R0 = 2



R0 = 2



Epidemics are Stochastic

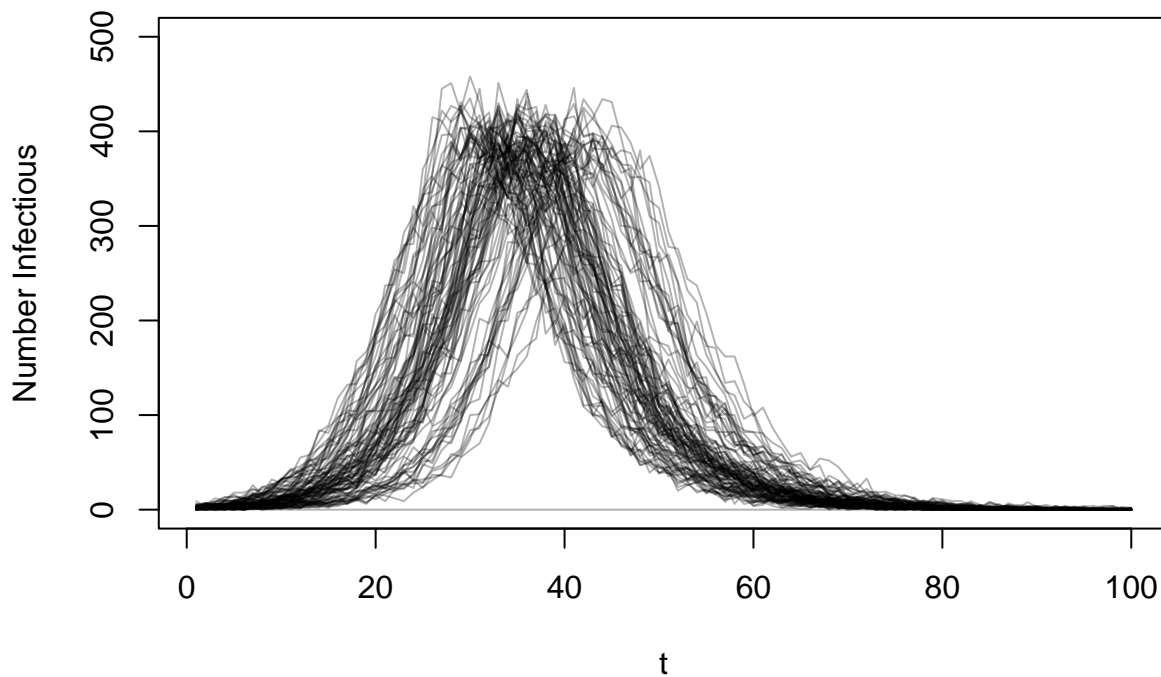
Here we simulate 100 epidemics from the same parameter values and plot the observed incidence curve from each simulation.

```
initsList <- list(beta = 0.4,
                  gamma = 0.2)
sirModel$setInits(initsList)

nSim <- 100

set.seed(1)
epiCurve <- matrix(NA, nrow = length(sirModel$Istar), ncol = nSim)
for (i in 1:nSim) {
  sirModel$simulate(nodesToSim, includeData = TRUE)
  epiCurve[,i] <- sirModel$Istar
}

plot(epiCurve[,1], type = 'l', col = adjustcolor('black', alpha = 0.3),
      ylim = c(0, 500), ylab = "Number Infectious", xlab='t')
for (i in 2:nSim) {
  lines(epiCurve[,i], col = adjustcolor('black', alpha = 0.3))
}
```



Model fitting to simulated data

Simulate data, then use it to fit the model.

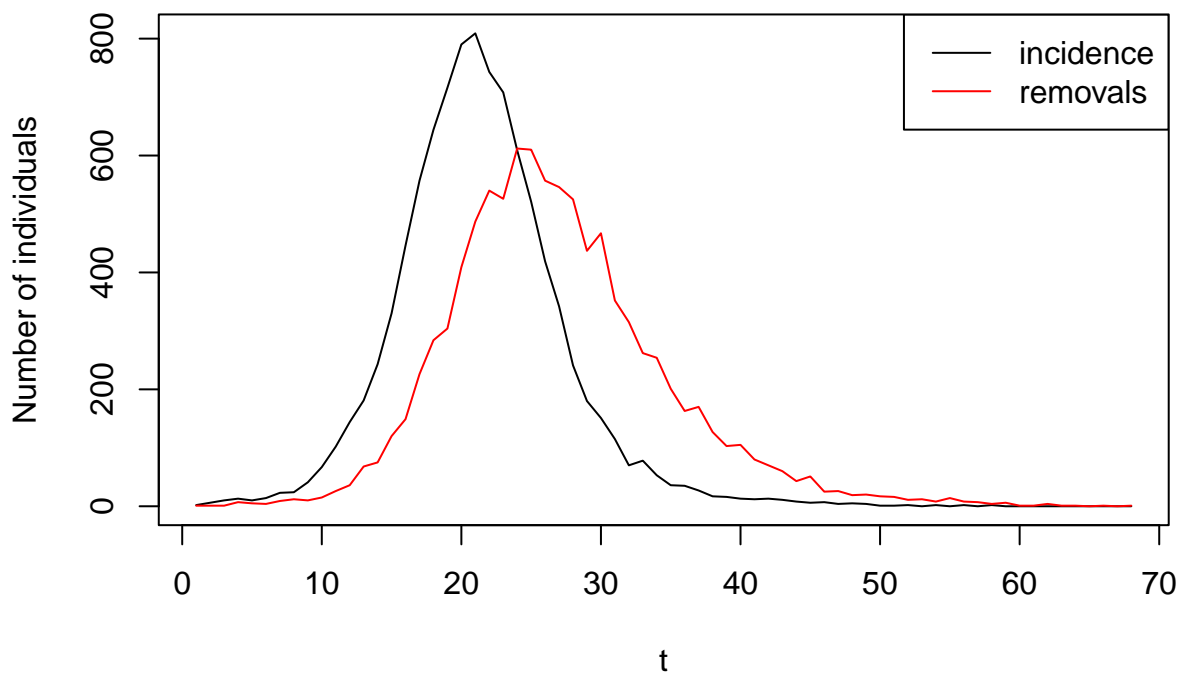
```
initsList <- list(beta = 0.6,
                 gamma = 0.2)
sirModel$setInits(initsList)

set.seed(1)
sirModel$simulate(nodesToSim, includeData = TRUE)

trueIstar <- sirModel$Istar
trueRstar <- sirModel$Rstar

endTime <- max(which(trueIstar > 0)) + 10
trueIstar <- trueIstar[1:endTime]
trueRstar <- trueRstar[1:endTime]

plot(trueIstar, type = 'l', ylab='Number of individuals', xlab='t')
lines(trueRstar, col = 'red')
legend('topright', c('incidence', 'removals'), col = c('black', 'red'), lwd = 1)
```



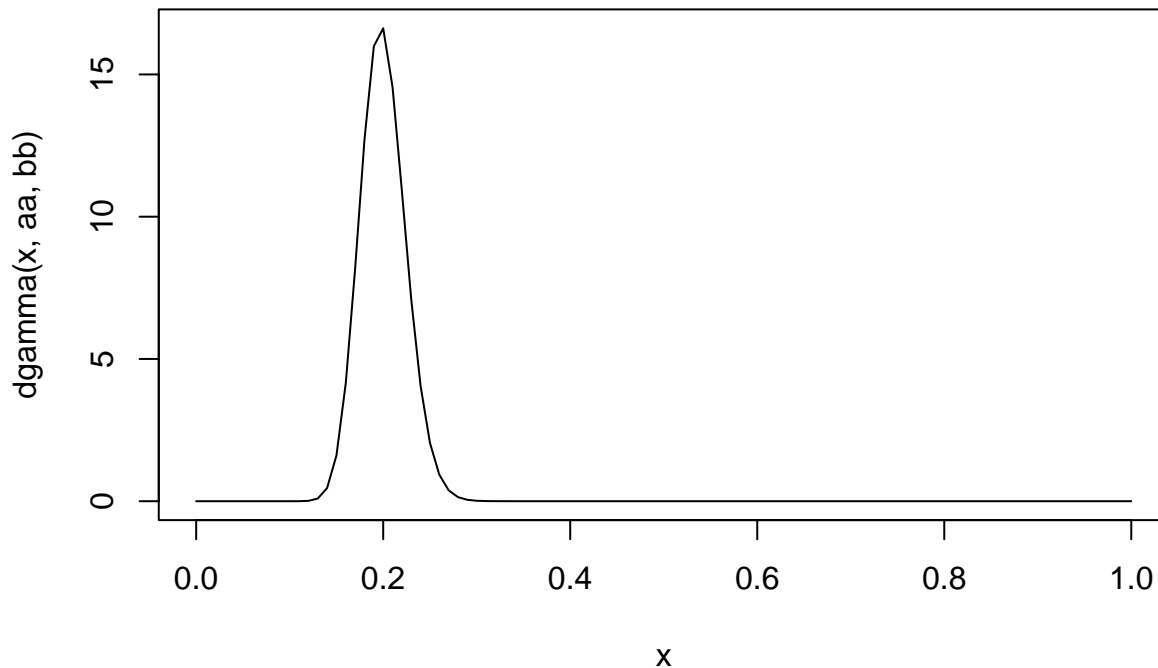
Model Specifications

Before fitting the model, we need to determine a reasonable prior for γ . The true value corresponds to a mean infectious period of 5 days, so we choose a prior that puts 90% probability on the mean infectious period between 4 and 6 days and is centered on 5 days.

```
bb <- 348
aa <- 0.2 * bb
pgamma(1/4, aa, bb) - pgamma(1/6, aa, bb)
```

```
## [1] 0.900183
```

```
curve(dgamma(x, aa, bb))
```



```
dataList <- list(Istar = trueIstar,
                Rstar = trueRstar)

constantsList <- list(N = 10000,
                    I0 = 5,
                    R0 = 0,
                    tau = length(dataList$Istar),
                    aa = aa,
                    bb = bb)

set.seed(2)
initsList <- list(beta = runif(1, 0, 1),
                 gamma = rgamma(1, aa, bb))

sirModelFit <- nimbleModel(SIR_code,
                        constants = constantsList,
                        data = dataList,
                        inits = initsList)
```

```
## Defining model
```

```

## Building model
## Setting data and initial values
## Running calculate on model
## [Note] Any error reports that follow may simply reflect missing values in model variables.
## Checking model sizes and dimensions
NIMBLE automatically calculates S, I, and R from Istar and Rstar, so these do not need to be inputs to the
model
with(sirModelFit, cbind(S, Istar, I, Rstar, R))[1:20,]

## Warning in cbind(S, Istar, I, Rstar, R): number of rows of result is not a
## multiple of vector length (arg 2)

##           S Istar   I Rstar   R
## [1,] 9995     2    5     1    0
## [2,] 9993     6    6     1    1
## [3,] 9987    10   11     1    2
## [4,] 9977    13   20     7    3
## [5,] 9964    10   26     5   10
## [6,] 9954    14   31     4   15
## [7,] 9940    23   41     9   19
## [8,] 9917    24   55    12   28
## [9,] 9893    41   67    10   40
## [10,] 9852    67   98    15   50
## [11,] 9785   102  150    26   65
## [12,] 9683   144  226    36   91
## [13,] 9539   181  334    68  127
## [14,] 9358   243  447    75  195
## [15,] 9115   330  615   120  270
## [16,] 8785   446  825   149  390
## [17,] 8339   557 1122   226  539
## [18,] 7782   644 1453   284  765
## [19,] 7138   716 1813   304 1049
## [20,] 6422   790 2225   409 1353

```


Use Default Configurations and Obtain Samples

Plotted with burn-in included here

```
myConfig <- configureMCMC(sirModelFit)
```

```
## ===== Monitors =====  
## thin = 1: beta, gamma  
## ===== Samplers =====  
## RW sampler (2)  
##   - beta  
##   - gamma
```

```
myMCMC <- buildMCMC(myConfig)
```

```
system.time({  
  compiled <- compileNimble(sirModelFit, myMCMC)  
  samples <- runMCMC(compiled$myMCMC, niter = 50000, setSeed = 3)  
})
```

```
## Compiling  
## [Note] This may take a minute.  
## [Note] Use 'showCompilerOutput = TRUE' to see C++ compilation details.
```

```
## Running chain 1 ...
```

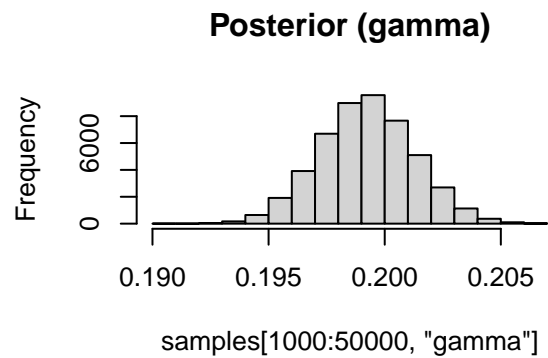
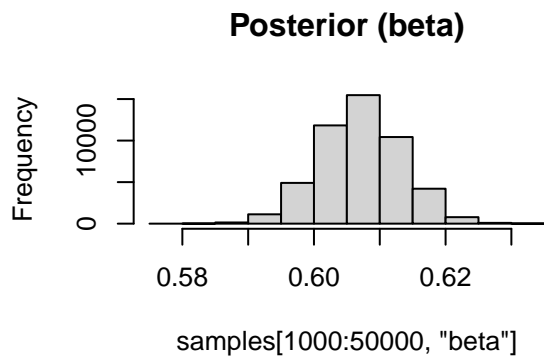
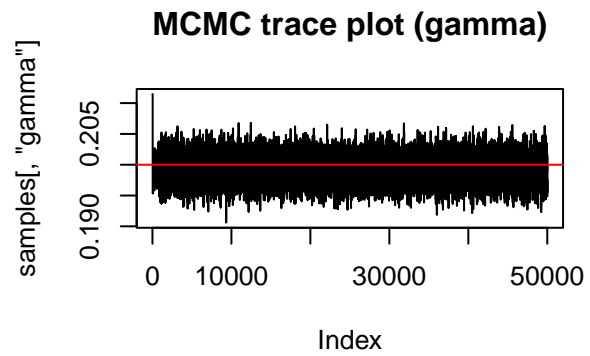
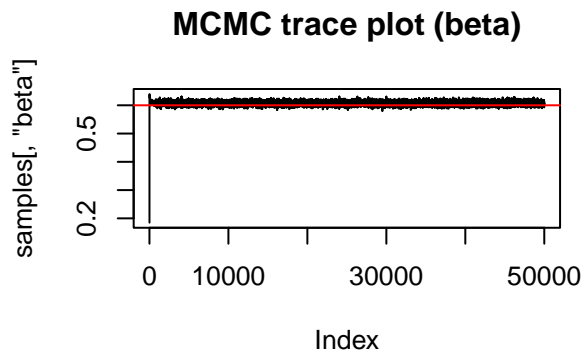
```
## |-----|-----|-----|-----|  
## |-----|-----|-----|-----|
```

```
##   user  system elapsed  
## 21.482   0.844  22.784
```

```
head(samples)
```

```
##           beta      gamma  
## [1,] 0.1848823 0.2114557  
## [2,] 0.4436705 0.2114557  
## [3,] 0.6394533 0.2114557  
## [4,] 0.6394533 0.2114557  
## [5,] 0.6394533 0.2114557  
## [6,] 0.6394533 0.2114557
```

```
par(mfrow = c(2,2))  
plot(samples[, 'beta'], type = 'l', main = 'MCMC trace plot (beta)')  
abline(h = 0.6, col = 'red')  
plot(samples[, 'gamma'], type = 'l', main = 'MCMC trace plot (gamma)')  
abline(h = 0.2, col = 'red')  
hist(samples[1000:50000, 'beta'], main = 'Posterior (beta)')  
hist(samples[1000:50000, 'gamma'], main = 'Posterior (gamma)')
```



```
#
# Posterior Mean and 95% Percentile Interval: Beta
mean(samples[1000:50000,'beta'])

## [1] 0.6070751
quantile(samples[1000:50000,'beta'], c(0.025,0.975))

##      2.5%      97.5%
## 0.5948491 0.6191422

#
# Posterior Mean and 95% Percentile Interval: Gamma
mean(samples[1000:50000,'gamma'])

## [1] 0.1992271
quantile(samples[1000:50000,'gamma'], c(0.025,0.975))

##      2.5%      97.5%
## 0.1953116 0.2032225

knitr::knit_exit()
```